

Predicting Wins in MLB

Author: Michael Tiernan

Brief History

Betting on sports has a long history despite the fact that it has only recently become legal to do so outside of Nevada (which legalized it in 1942). New Jersey became just the second state in the US to legalize sports betting following a US Supreme Court ruling in May of 2018. This was a major court decision that opened the floodgates for more states to follow and billions of dollars to be generated doing so. As of today, June 24, 2021, 22 states have legalized sports betting in the US. The rest of the US has either passed or introduced a bill to legalize the industry with the exception of only three states (Idaho, Utah, Wisconsin) who have not.

For those not familiar with the industry, a sportsbook (ex. DraftKings, William Hill, FanDuel, Barstool) creates betting odds based on their own machine learning models for the user to then wager money on. If the two teams are very similar and there is no clear favorite, the odds for that game would be close to even. When the odds are even, a 100betwillreturna100 payout (if the team you bet on wins). But what if one team is heavily favored? When the best team has a game against the worst team, the odds for that game have to account for this disparity and the way they do that is adjusting the odds to let's say -200, which means that your 100betnowonlyreturnsa50 payout. By doing so, one may be persuaded to then bet on the underdog in this scenario because the odds on the underdog will be +150, which means your 100wouldreturn150 payout. By adjusting the odds, the sportsbook is looking to have an even amount of action on each team because this almost assures them a profit due to the fact that it is expected that the favorite will win, in which case they only have to payout 50inthe2ndscenario, butatthesametimecollectingthe 100 lost bet from the underdog bettor. This, coupled with the resources and mastery that the sportsbooks have in creating the odds, is what makes betting on sports so difficult in terms of being profitable over the long-term for the average user.

Business Problem

Given this brief history, it is clear that the average bettor has an uphill battle to long-term profitability. For my analysis, I am looking to answer the two questions below:

- 1) What are the best predictors for season total wins?
- 2) What are the best predictors for individual game wins?

Like any other game, whether it be casino, board games, or competitive sport, the intelligent player is going to look for a competitive advantage. Within the sports betting industry, the inexperienced bettor likely makes their wagers based off of gut intuition. The intermediate bettor may incorporate

research from articles read or a comparison of some surface-level stats, but the true competitive advantage comes from a deep-dive into the data, using machine learning algorithms to account for as many factors/predictors as possible.

This is the strategy I will continuously build upon in pursuit of creating a highly competitive algorithm for long-term sports betting profitability.

```
In [1]: # importing packages for analysis
import pandas as pd
# setting pandas display to avoid scientific notation
pd.options.display.float_format = '{:.3f}'.format
# setting to see all columns in outputs of dataframes
pd.set_option('display.max_columns', None)
import numpy as np
import matplotlib.pyplot as plt
%matplotlib inline
import seaborn as sns
sns.set_theme(style="whitegrid")
import warnings
warnings.filterwarnings('ignore')
```

Yearly Stats

This first dataset comes from the lahman's baseball database. This is the go-to database for baseball analysis. This particular dataset that I am going to work with has yearly data for each team from 1871 - 2020. There are 48 columns of stats and information to get me started.

The Covid-19 pandemic caused the 2020 season to be shortened to just 60 games (a regular full season is 162 games) and some players opting out of playing. Given this circumstance, I chose to remove this shortened year from my analysis.

```
In [2]: df_teams = pd.read_csv('data/Teams.csv')
df_teams.tail()
```

Out[2]:

| | yearID | lgID | teamID | franchID | divID | Rank | G | Ghome | W | L | DivWin | WCWin | LgWin | V |
|-------------|--------|------|--------|----------|-------|------|----|--------|----|----|--------|-------|-------|---|
| 2950 | 2020 | NL | SLN | STL | C | 3 | 58 | 27.000 | 30 | 28 | N | Y | N | |
| 2951 | 2020 | AL | TBA | TBD | E | 1 | 60 | 29.000 | 40 | 20 | Y | N | Y | |
| 2952 | 2020 | AL | TEX | TEX | W | 5 | 60 | 30.000 | 22 | 38 | N | N | N | |
| 2953 | 2020 | AL | TOR | TOR | E | 3 | 60 | 26.000 | 32 | 28 | N | Y | N | |
| 2954 | 2020 | NL | WAS | WSN | E | 4 | 60 | 33.000 | 26 | 34 | N | N | N | |

```
In [3]: # drop cols
cols = ['teamID', 'lgID', 'divID', 'Rank', 'G', 'Ghome', 'L', 'CS', 'HBP', 'IPouts', 'name']
df_teams.drop(cols, axis=1, inplace=True)
df_teams.tail()
```

Out[3]:

| | yearID | franchID | W | DivWin | WCWin | LgWin | WSWin | R | AB | H | 2B | 3B | HR | B |
|-------------|--------|----------|----|--------|-------|-------|-------|-----|------|-----|-----|----|----|--------|
| 2950 | 2020 | STL | 30 | N | Y | N | N | 240 | 1752 | 410 | 73 | 7 | 51 | 205.00 |
| 2951 | 2020 | TBD | 40 | Y | N | Y | N | 289 | 1975 | 470 | 105 | 12 | 80 | 243.00 |
| 2952 | 2020 | TEX | 22 | N | N | N | N | 224 | 1936 | 420 | 80 | 9 | 62 | 167.00 |
| 2953 | 2020 | TOR | 32 | N | Y | N | N | 302 | 2023 | 516 | 104 | 4 | 88 | 203.00 |
| 2954 | 2020 | WSN | 26 | N | N | N | N | 293 | 1968 | 519 | 112 | 12 | 66 | 192.00 |

```
In [4]: # dummy championships
cols = ['DivWin', 'WCWin', 'LgWin', 'WSWin']
dummy_df = pd.get_dummies(df_teams[cols], drop_first=True)
dummy_df.tail()
```

Out[4]:

| | DivWin_Y | WCWin_Y | LgWin_Y | WSWin_Y |
|-------------|----------|---------|---------|---------|
| 2950 | 0 | 1 | 0 | 0 |
| 2951 | 1 | 0 | 1 | 0 |
| 2952 | 0 | 0 | 0 | 0 |
| 2953 | 0 | 1 | 0 | 0 |
| 2954 | 0 | 0 | 0 | 0 |

```
In [5]: df_teams = pd.concat([df_teams, dummy_df], axis=1)
df_teams.drop(['DivWin', 'WCWin', 'LgWin', 'WSWin'], axis=1, inplace=True)
df_teams.tail()
```

Out[5]:

| | yearID | franchID | W | R | AB | H | 2B | 3B | HR | BB | SO | SB | SF | RA |
|-------------|--------|----------|----|-----|------|-----|-----|----|----|---------|---------|--------|--------|-----|
| 2950 | 2020 | STL | 30 | 240 | 1752 | 410 | 73 | 7 | 51 | 205.000 | 477.000 | 18.000 | 16.000 | 229 |
| 2951 | 2020 | TBD | 40 | 289 | 1975 | 470 | 105 | 12 | 80 | 243.000 | 608.000 | 48.000 | 14.000 | 229 |
| 2952 | 2020 | TEX | 22 | 224 | 1936 | 420 | 80 | 9 | 62 | 167.000 | 548.000 | 49.000 | 18.000 | 312 |
| 2953 | 2020 | TOR | 32 | 302 | 2023 | 516 | 104 | 4 | 88 | 203.000 | 508.000 | 33.000 | 14.000 | 312 |
| 2954 | 2020 | WSN | 26 | 293 | 1968 | 519 | 112 | 12 | 66 | 192.000 | 451.000 | 33.000 | 21.000 | 301 |

```
In [6]: # feature engineer team run diff and batting avg.
df_teams['run_diff'] = (df_teams['R'] - df_teams['RA'])
df_teams['ba'] = (df_teams['H'] / df_teams['AB'])
df_teams.tail()
```

Out[6]:

| | yearID | franchID | W | R | AB | H | 2B | 3B | HR | BB | SO | SB | SF | RA |
|-------------|--------|----------|----|-----|------|-----|-----|----|----|---------|---------|--------|--------|-----|
| 2950 | 2020 | STL | 30 | 240 | 1752 | 410 | 73 | 7 | 51 | 205.000 | 477.000 | 18.000 | 16.000 | 229 |
| 2951 | 2020 | TBD | 40 | 289 | 1975 | 470 | 105 | 12 | 80 | 243.000 | 608.000 | 48.000 | 14.000 | 229 |
| 2952 | 2020 | TEX | 22 | 224 | 1936 | 420 | 80 | 9 | 62 | 167.000 | 548.000 | 49.000 | 18.000 | 312 |
| 2953 | 2020 | TOR | 32 | 302 | 2023 | 516 | 104 | 4 | 88 | 203.000 | 508.000 | 33.000 | 14.000 | 312 |
| 2954 | 2020 | WSN | 26 | 293 | 1968 | 519 | 112 | 12 | 66 | 192.000 | 451.000 | 33.000 | 21.000 | 301 |

```
In [7]: # filter db to years 2000 - 2019
df_teams = df_teams.loc[(df_teams['yearID']>1999) & (df_teams['yearID']<2020)]
df_teams
```

Out[7]:

| | yearID | franchID | W | R | AB | H | 2B | 3B | HR | BB | SO | SB | SF |
|-------------|--------|----------|-----|-----|------|------|-----|-----|-----|---------|----------|---------|--------|
| 2325 | 2000 | ANA | 82 | 864 | 5628 | 1574 | 309 | 34 | 236 | 608.000 | 1024.000 | 93.000 | 43.000 |
| 2326 | 2000 | ARI | 85 | 792 | 5527 | 1466 | 282 | 44 | 179 | 535.000 | 975.000 | 97.000 | 58.000 |
| 2327 | 2000 | ATL | 95 | 810 | 5489 | 1490 | 274 | 26 | 179 | 595.000 | 1010.000 | 148.000 | 45.000 |
| 2328 | 2000 | BAL | 74 | 794 | 5549 | 1508 | 310 | 22 | 184 | 558.000 | 900.000 | 126.000 | 54.000 |
| 2329 | 2000 | BOS | 85 | 792 | 5630 | 1503 | 316 | 32 | 167 | 611.000 | 1019.000 | 43.000 | 48.000 |
| ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... |
| 2920 | 2019 | STL | 91 | 764 | 5449 | 1336 | 246 | 24 | 210 | 561.000 | 1420.000 | 116.000 | 39.000 |
| 2921 | 2019 | TBD | 96 | 769 | 5628 | 1427 | 291 | 29 | 217 | 542.000 | 1493.000 | 94.000 | 34.000 |
| 2922 | 2019 | TEX | 78 | 810 | 5540 | 1374 | 296 | 24 | 223 | 534.000 | 1578.000 | 131.000 | 44.000 |
| 2923 | 2019 | TOR | 67 | 726 | 5493 | 1299 | 270 | 21 | 247 | 509.000 | 1514.000 | 51.000 | 28.000 |
| 2924 | 2019 | WSN | 93 | 873 | 5512 | 1460 | 298 | 27 | 231 | 584.000 | 1308.000 | 116.000 | 42.000 |

```
In [8]: nym = df_teams.loc[df_teams['franchID']=='NYM']
        nym
```

Out[8]:

| | yearID | franchID | W | R | AB | H | 2B | 3B | HR | BB | SO | SB | SF | R |
|-------------|--------|----------|----|-----|------|------|-----|----|-----|---------|----------|---------|--------|----|
| 2344 | 2000 | NYM | 94 | 807 | 5486 | 1445 | 281 | 20 | 198 | 675.000 | 1037.000 | 66.000 | 51.000 | 70 |
| 2374 | 2001 | NYM | 82 | 642 | 5459 | 1361 | 273 | 18 | 147 | 545.000 | 1062.000 | 66.000 | 35.000 | 70 |
| 2404 | 2002 | NYM | 75 | 690 | 5496 | 1409 | 238 | 22 | 160 | 486.000 | 1044.000 | 87.000 | 30.000 | 70 |
| 2434 | 2003 | NYM | 66 | 642 | 5341 | 1317 | 262 | 24 | 124 | 489.000 | 1035.000 | 70.000 | 45.000 | 70 |
| 2464 | 2004 | NYM | 71 | 684 | 5532 | 1376 | 289 | 20 | 185 | 512.000 | 1159.000 | 107.000 | 34.000 | 70 |
| 2493 | 2005 | NYM | 83 | 722 | 5505 | 1421 | 279 | 32 | 175 | 486.000 | 1075.000 | 153.000 | 38.000 | 64 |
| 2523 | 2006 | NYM | 97 | 834 | 5558 | 1469 | 323 | 41 | 200 | 547.000 | 1071.000 | 146.000 | 47.000 | 70 |
| 2553 | 2007 | NYM | 88 | 804 | 5605 | 1543 | 294 | 27 | 177 | 549.000 | 981.000 | 200.000 | 58.000 | 70 |
| 2583 | 2008 | NYM | 89 | 799 | 5606 | 1491 | 274 | 38 | 172 | 619.000 | 1024.000 | 138.000 | 49.000 | 70 |
| 2613 | 2009 | NYM | 70 | 671 | 5453 | 1472 | 295 | 49 | 95 | 526.000 | 928.000 | 122.000 | 55.000 | 70 |
| 2643 | 2010 | NYM | 79 | 656 | 5465 | 1361 | 266 | 40 | 128 | 502.000 | 1095.000 | 130.000 | 57.000 | 60 |
| 2673 | 2011 | NYM | 77 | 718 | 5600 | 1477 | 309 | 39 | 108 | 571.000 | 1085.000 | 130.000 | 48.000 | 70 |
| 2703 | 2012 | NYM | 74 | 650 | 5450 | 1357 | 286 | 21 | 139 | 503.000 | 1250.000 | 79.000 | 30.000 | 70 |
| 2733 | 2013 | NYM | 74 | 619 | 5559 | 1318 | 263 | 32 | 130 | 512.000 | 1384.000 | 114.000 | 32.000 | 60 |
| 2763 | 2014 | NYM | 79 | 629 | 5472 | 1306 | 275 | 19 | 125 | 516.000 | 1264.000 | 101.000 | 44.000 | 60 |
| 2793 | 2015 | NYM | 90 | 683 | 5527 | 1351 | 295 | 17 | 177 | 488.000 | 1290.000 | 51.000 | 32.000 | 60 |
| 2823 | 2016 | NYM | 87 | 671 | 5459 | 1342 | 240 | 19 | 218 | 517.000 | 1302.000 | 42.000 | 41.000 | 60 |
| 2853 | 2017 | NYM | 70 | 735 | 5510 | 1379 | 286 | 28 | 224 | 529.000 | 1291.000 | 58.000 | 37.000 | 80 |
| 2883 | 2018 | NYM | 77 | 676 | 5468 | 1282 | 265 | 34 | 170 | 566.000 | 1404.000 | 71.000 | 42.000 | 70 |
| 2913 | 2019 | NYM | 86 | 791 | 5624 | 1445 | 280 | 17 | 242 | 516.000 | 1384.000 | 56.000 | 27.000 | 70 |

Which team has averaged the most amount of wins over the past 20 years?

As a Mets fan, I begrudgingly hypothesize that the Yankees is our answer.


```
In [9]: wins_df = df_teams.groupby(['franchID'])['W'].mean()
```

```
In [10]: wins_df = pd.DataFrame(wins_df).sort_values('W', ascending=False)
wins_df
```

Out[10]:

| | W |
|-----------------|--------|
| franchID | |
| NYN | 94.300 |
| STL | 90.600 |
| BOS | 89.600 |
| LAD | 89.050 |
| ATL | 86.750 |
| OAK | 86.450 |
| ANA | 86.100 |
| SFG | 83.800 |
| CLE | 83.550 |
| PHI | 81.850 |

```
In [13]: fig,ax = plt.subplots(figsize=(16,10))
ax = sns.barplot(x=wins_df.index, y='W', data=wins_df, palette='Blues_d')
ax.set_title("Average Regular Season Wins: P20Years", fontsize=18)
ax.set(xlabel='Franchise', ylabel='Wins')
plt.show()
```



What about past 5 years?

```
In [15]: # filter db to P5Years, 2015 - 2019
df_p5y = df_teams.loc[(df_teams['yearID']>2014)]
df_p5y
```

Out[15]:

| | yearID | franchID | W | R | AB | H | 2B | 3B | HR | BB | SO | SB | SF | R |
|-------------|--------|----------|-----|-----|------|------|-----|-----|-----|---------|----------|---------|--------|-----|
| 2775 | 2015 | ARI | 79 | 720 | 5649 | 1494 | 289 | 48 | 154 | 490.000 | 1312.000 | 132.000 | 57.000 | 7 |
| 2776 | 2015 | ATL | 67 | 573 | 5420 | 1361 | 251 | 18 | 100 | 471.000 | 1107.000 | 69.000 | 31.000 | 76 |
| 2777 | 2015 | BAL | 81 | 713 | 5485 | 1370 | 246 | 20 | 217 | 418.000 | 1331.000 | 44.000 | 32.000 | 69 |
| 2778 | 2015 | BOS | 78 | 748 | 5640 | 1495 | 294 | 33 | 161 | 478.000 | 1148.000 | 71.000 | 42.000 | 75 |
| 2779 | 2015 | CHW | 76 | 622 | 5533 | 1381 | 260 | 27 | 136 | 404.000 | 1231.000 | 68.000 | 37.000 | 70 |
| ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... |
| 2920 | 2019 | STL | 91 | 764 | 5449 | 1336 | 246 | 24 | 210 | 561.000 | 1420.000 | 116.000 | 39.000 | 66 |
| 2921 | 2019 | TBD | 96 | 769 | 5628 | 1427 | 291 | 29 | 217 | 542.000 | 1493.000 | 94.000 | 34.000 | 65 |
| 2922 | 2019 | TEX | 78 | 810 | 5540 | 1374 | 296 | 24 | 223 | 534.000 | 1578.000 | 131.000 | 44.000 | 87 |
| 2923 | 2019 | TOR | 67 | 726 | 5493 | 1299 | 270 | 21 | 247 | 509.000 | 1514.000 | 51.000 | 28.000 | 82 |
| 2924 | 2019 | WSN | 93 | 873 | 5512 | 1460 | 298 | 27 | 231 | 584.000 | 1308.000 | 116.000 | 42.000 | 72 |

150 rows × 34 columns

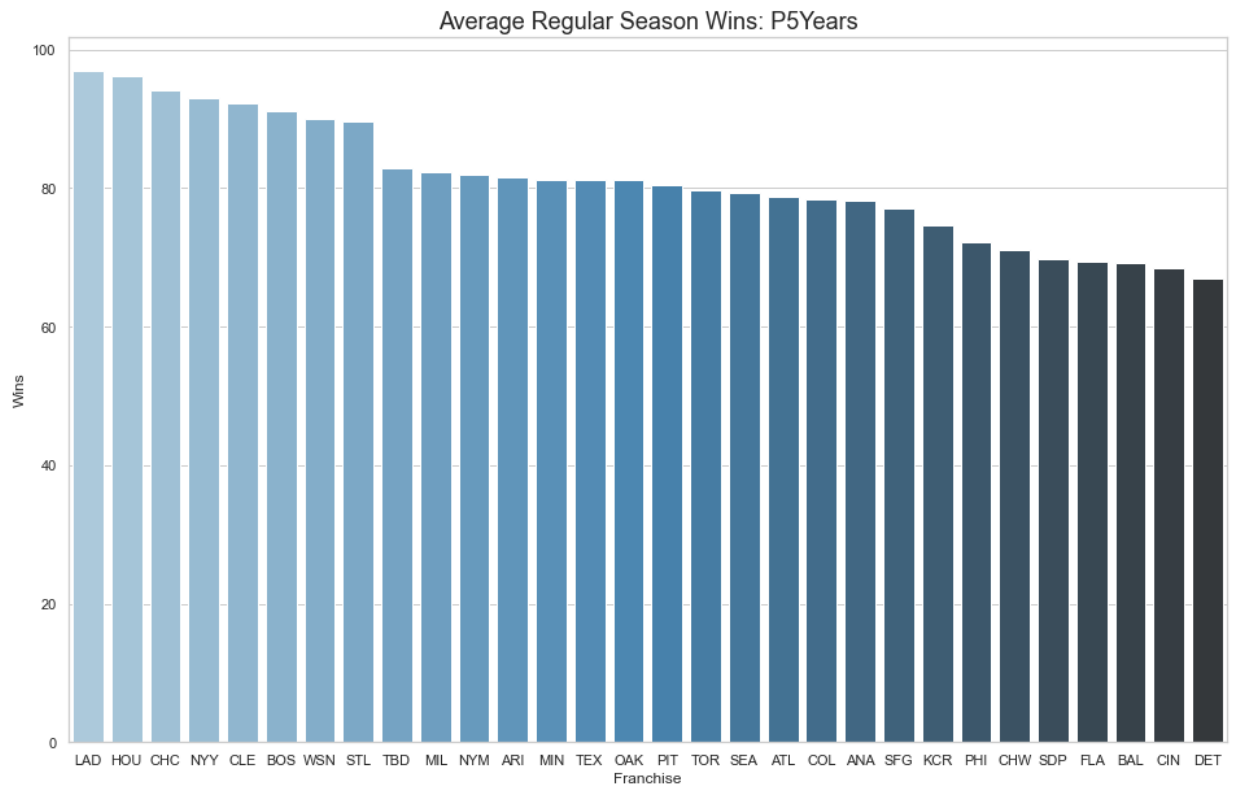
```
In [17]: df_p5y_grouped = df_p5y.groupby(['franchID'])['W'].mean()
```

```
In [18]: df_p5y_grouped = pd.DataFrame(df_p5y_grouped).sort_values('W', ascending=False)
df_p5y_grouped
```

Out[18]:

| | W |
|------------|--------|
| franchID | |
| LAD | 97.000 |
| HOU | 96.200 |
| CHC | 94.200 |
| NYN | 93.000 |
| CLE | 92.200 |
| BOS | 91.200 |
| WSN | 90.000 |
| STL | 89.600 |
| TBD | 82.800 |
| MIL | 82.400 |

```
In [19]: # P5Years
fig,ax = plt.subplots(figsize=(16,10))
ax = sns.barplot(x=df_p5y_grouped.index, y='W', data=df_p5y_grouped, palette='Blu
ax.set_title("Average Regular Season Wins: P5Years", fontsize=18)
ax.set(xlabel='Franchise', ylabel='Wins')
plt.show()
```



```
In [20]: # P20Years and P5Years dfs
print(df_teams.shape)
print(df_p5y.shape)
```

```
(600, 34)
```


```
(150, 34)
```


In [21]: df_p5y

Out[21]:

| | yearID | franchID | W | R | AB | H | 2B | 3B | HR | BB | ... | DP | FP | BPF | PPF | Div' |
|-------------|--------|----------|-----|-----|------|------|-----|-----|-----|-------|-----|-----|-------|-----|-----|------|
| 2775 | 2015 | ARI | 79 | 720 | 5649 | 1494 | 289 | 48 | 154 | 490.0 | ... | 146 | 0.986 | 107 | 106 | |
| 2776 | 2015 | ATL | 67 | 573 | 5420 | 1361 | 251 | 18 | 100 | 471.0 | ... | 186 | 0.985 | 97 | 97 | |
| 2777 | 2015 | BAL | 81 | 713 | 5485 | 1370 | 246 | 20 | 217 | 418.0 | ... | 134 | 0.987 | 103 | 104 | |
| 2778 | 2015 | BOS | 78 | 748 | 5640 | 1495 | 294 | 33 | 161 | 478.0 | ... | 148 | 0.984 | 104 | 107 | |
| 2779 | 2015 | CHW | 76 | 622 | 5533 | 1381 | 260 | 27 | 136 | 404.0 | ... | 159 | 0.983 | 92 | 93 | |
| ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | |
| 2920 | 2019 | STL | 91 | 764 | 5449 | 1336 | 246 | 24 | 210 | 561.0 | ... | 168 | 0.989 | 98 | 97 | |
| 2921 | 2019 | TBD | 96 | 769 | 5628 | 1427 | 291 | 29 | 217 | 542.0 | ... | 126 | 0.985 | 97 | 96 | |
| 2922 | 2019 | TEX | 78 | 810 | 5540 | 1374 | 296 | 24 | 223 | 534.0 | ... | 143 | 0.982 | 111 | 112 | |
| 2923 | 2019 | TOR | 67 | 726 | 5493 | 1299 | 270 | 21 | 247 | 509.0 | ... | 141 | 0.984 | 97 | 98 | |
| 2924 | 2019 | WSN | 93 | 873 | 5512 | 1460 | 298 | 27 | 231 | 584.0 | ... | 111 | 0.985 | 106 | 104 | |

150 rows × 34 columns



Advanced Stats

Baseballsavant.com allows you to dynamically pull stats. For my analysis I'm going to use linedrivepercentage, xOBP, and xwOBA.

- linedrivepercentage = percentage of time a team hit a linedrive. Linedrives are a hard-hit ball that has the best opportunity of the batter to successfully reach base, as opposed to groundballs and flyballs.
- xOBP = expected on base percentage, which is formulated by quality of contact, strikeouts, and walks
- xwOBA = expected weighted on base average, which is formulated using exit velocity, launch angle, and on certain types of batted balls, sprint speed

```
In [21]: df_newstats = pd.read_csv('data/new_stats.csv')
df_newstats
```

Out[21]:

| | franchID | yearID | linedrivepercent | xobp | xwoba |
|-----|----------|--------|------------------|-------|-------|
| 0 | ANA | 2015 | 27.200 | 0.304 | 0.304 |
| 1 | ARI | 2015 | 29.100 | 0.306 | 0.303 |
| 2 | ATL | 2015 | 25.800 | 0.307 | 0.294 |
| 3 | BAL | 2015 | 26.400 | 0.297 | 0.307 |
| 4 | BOS | 2015 | 23.900 | 0.307 | 0.304 |
| ... | ... | ... | ... | ... | ... |
| 145 | STL | 2019 | 24.900 | 0.317 | 0.315 |
| 146 | TBD | 2019 | 25.600 | 0.326 | 0.327 |
| 147 | TEX | 2019 | 25.600 | 0.315 | 0.317 |
| 148 | TOR | 2019 | 24.500 | 0.307 | 0.312 |
| 149 | WSN | 2019 | 25.100 | 0.334 | 0.334 |

150 rows × 5 columns

```
In [46]: xwoba = df_newstats.groupby(['franchID', 'yearID'])['xwoba'].mean()
```

```
In [66]: xwoba_df = pd.DataFrame(xwoba).sort_values(['yearID', 'xwoba'], ascending=False)
xwoba_df.head(10)
```

Out[66]:

| | | xwoba |
|----------|--------|-------|
| franchID | yearID | |
| MIN | 2019 | 0.345 |
| ATL | 2019 | 0.336 |
| WSN | 2019 | 0.334 |
| LAD | 2019 | 0.333 |
| NYG | 2019 | 0.333 |
| HOU | 2019 | 0.332 |
| BOS | 2019 | 0.331 |
| OAK | 2019 | 0.331 |
| TBD | 2019 | 0.327 |
| CHC | 2019 | 0.324 |

The Minnesota Twins record in 2019 was 101-61

The Boston Red Sox record in 2018 was 108-54


You can clearly see how these advanced stats are great predictors for team performance. Unfortunately, these advanced stats only go back to the year 2015, so they will not be used in my models at this time.

```
In [67]: xwoba_df.reset_index(inplace=True)
```

```
In [72]: plt.figure(figsize=(32,14))
```

```
plt.subplot(321)
sns.lineplot(xwoba_df.franchID, xwoba_df.xwoba, hue=xwoba_df.yearID, palette='Blu')
plt.title('Team xwOBA By Year', fontsize=14)
```

```
Out[72]: Text(0.5, 1.0, 'Team xwOBA By Year')
```



Vegas Lines

I was able to find the lines set by the bookmakers for MLB season total wins for each team over the past 6 years. I was excited to be able to find this because this allows me to compare the vegas lines vs. the actual outcomes. As you can imagine, this is not something that is readily available by a simple google search because this is like finding the data for a casino and knowing exactly how accurate/profitable each game and table really is. Lets dig in!

```
In [254]: vegas_odds = pd.read_csv('data/vegas_mlb_seaswins.csv')
```

```
In [255]: vegas_odds.drop_duplicates(inplace=True)
```

```
In [256]: vegas_odds.tail()
```

```
Out[256]:
```

| | year | Team | vegas_win_total |
|-----|----------|------|-----------------|
| 176 | 2019.000 | TBD | 85.500 |
| 177 | 2019.000 | TEX | 70.500 |
| 178 | 2019.000 | TOR | 76.500 |
| 179 | 2019.000 | WSN | 88.500 |
| 180 | nan | NaN | nan |

```
In [257]: # remove last row - should be 180 rows
vegas_odds = vegas_odds.iloc[:-1 , :]
```

```
In [259]: # rename columns and lowercase all columns
vegas_odds = vegas_odds.rename(columns={'year': 'yearID',
                                         'Team': 'franchID'})
```

```
In [261]: vegas_odds.set_index(['yearID', 'franchID'], inplace=True)
```

```
In [116]: df_teams.set_index(['yearID', 'franchID'], inplace=True)
```

...

```
In [262]: # join the vegas odds to the main teams df
vegas_df = df_teams.join(vegas_odds, how='inner', lsuffix='teams_', rsuffix='stat')
vegas_df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
MultiIndex: 180 entries, (2017.0, LAD) to (2016.0, WSN)
Data columns (total 33 columns):
W                180 non-null int64
R                180 non-null int64
AB               180 non-null int64
H                180 non-null int64
2B               180 non-null int64
3B               180 non-null int64
HR               180 non-null int64
BB               180 non-null float64
SO               180 non-null float64
SB               180 non-null float64
SF               180 non-null float64
RA               180 non-null int64
ER               180 non-null int64
ERA              180 non-null float64
CG               180 non-null int64
SHO              180 non-null int64
SV               180 non-null int64
HA               180 non-null int64
HRA              180 non-null int64
BBA              180 non-null int64
SOA              180 non-null int64
E                180 non-null int64
DP               180 non-null int64
FP               180 non-null float64
BPF              180 non-null int64
PPF              180 non-null int64
DivWin_Y         180 non-null uint8
WCWin_Y           180 non-null uint8
LgWin_Y           180 non-null uint8
WSWin_Y           180 non-null uint8
run_diff          180 non-null int64
ba                180 non-null float64
vegas_win_total   180 non-null float64
dtypes: float64(8), int64(21), uint8(4)
memory usage: 42.3+ KB
```

```
In [264]: vegas_df.reset_index(inplace=True)
```

```
In [265]: vegas_df['yearID'] = pd.to_datetime(vegas_df['yearID'], format='%Y')
```


```
In [267]: vegas_df_2019 = vegas_df.loc[vegas_df['yearID']=='2019-01-01']
vegas_df_2018 = vegas_df.loc[vegas_df['yearID']=='2018-01-01']
vegas_df_2017 = vegas_df.loc[vegas_df['yearID']=='2017-01-01']
vegas_df_2016 = vegas_df.loc[vegas_df['yearID']=='2016-01-01']
vegas_df_2015 = vegas_df.loc[vegas_df['yearID']=='2015-01-01']
vegas_df_2014 = vegas_df.loc[vegas_df['yearID']=='2014-01-01']
```

```
In [268]: plt.figure(figsize=(20,8))
X = vegas_df_2019['franchID']
y = vegas_df_2019['W']
z = vegas_df_2019['vegas_win_total']

plt.plot(X, y, color='navy', label='Actual Wins')
plt.plot(X, z, color='red', label='Vegas Line')

plt.xlabel('Teams', fontsize=14)
plt.ylabel('Season Total Wins', fontsize=14)
plt.title('2019 Season Total Wins: Actual vs. Vegas Line', fontsize=18)

plt.legend(fontsize='large')
plt.show()
```



```
In [322]: fig, (ax1, ax2, ax3, ax4, ax5, ax6) = plt.subplots(6, sharex=True, sharey=True,
fig.suptitle('Season Total Wins: Actual vs. Vegas Line', fontsize=22)
fig.text(0.5,0.04, "Teams", ha="center", va="center", fontsize=14)
fig.text(0.05,0.5, "Season Total Wins", ha="center", va="center", rotation=90, fo

X = vegas_df_2019['franchID']
y = vegas_df_2019['W']
z = vegas_df_2019['vegas_win_total']

w = vegas_df_2018['W']
b = vegas_df_2018['vegas_win_total']

v = vegas_df_2017['W']
c = vegas_df_2017['vegas_win_total']

u = vegas_df_2016['W']
d = vegas_df_2016['vegas_win_total']

t = vegas_df_2015['W']
e = vegas_df_2015['vegas_win_total']

s = vegas_df_2014['W']
f = vegas_df_2014['vegas_win_total']

ax1.plot(X, y, color='navy', label='Actual Wins')
ax1.plot(X, z, color='red', label='Vegas Line')

ax2.plot(X, w, color='navy', label='Actual Wins')
ax2.plot(X, b, color='red', label='Vegas Line')

ax3.plot(X, v, color='navy', label='Actual Wins')
ax3.plot(X, c, color='red', label='Vegas Line')

ax4.plot(X, u, color='navy', label='Actual Wins')
ax4.plot(X, d, color='red', label='Vegas Line')

ax5.plot(X, t, color='navy', label='Actual Wins')
ax5.plot(X, e, color='red', label='Vegas Line')


ax6.plot(X, s, color='navy', label='Actual Wins')
ax6.plot(X, f, color='red', label='Vegas Line')

ax1.set_title('2019', fontsize=14)
ax2.set_title('2018', fontsize=14)
ax3.set_title('2017', fontsize=14)
ax4.set_title('2016', fontsize=14)
ax5.set_title('2015', fontsize=14)
ax6.set_title('2014', fontsize=14)

ax1.legend(fontsize='medium', loc='upper left')
plt.savefig('./images/vegas_accuracy.png')
plt.show();
```

Preprocess_and_EDA

Season Total Wins: Actual vs. Vegas Line



```
In [323]: # calculate absolute diff wins - vegas line
# calculate percent (diff wins - vegas line)/vegas line
vegas_df['abs_line_accuracy'] = vegas_df['W'] - vegas_df['vegas_win_total']
vegas_df['pdiff_line_accuracy'] = (vegas_df['W'] - vegas_df['vegas_win_total']) /
```



```
In [324]: # Vegas accuracy by year and team 2014 - 2019
compare_lines = vegas_df.groupby(['yearID', 'franchID'])['W', 'vegas_win_total', 'abs_line_acc']
compare_lines
```

Out[324]:

| | | W | vegas_win_total | abs_line_accuracy | pdiff_line_accuracy |
|------------|----------|-----|-----------------|-------------------|---------------------|
| yearID | franchID | | | | |
| 2014-01-01 | ANA | 98 | 87.000 | 11.000 | 0.126 |
| | ARI | 64 | 80.500 | -16.500 | -0.205 |
| | ATL | 79 | 87.500 | -8.500 | -0.097 |
| | BAL | 96 | 81.500 | 14.500 | 0.178 |
| | BOS | 71 | 88.000 | -17.000 | -0.193 |
| ... | ... | ... | ... | ... | ... |
| 2019-01-01 | STL | 91 | 88.500 | 2.500 | 0.028 |
| | TBD | 96 | 85.500 | 10.500 | 0.123 |
| | TEX | 78 | 70.500 | 7.500 | 0.106 |
| | TOR | 67 | 76.500 | -9.500 | -0.124 |
| | WSN | 93 | 88.500 | 4.500 | 0.051 |

180 rows × 4 columns

```
In [329]: # Vegas accuracy by year cumulative of all teams
compare_lines2 = vegas_df.groupby(['yearID'])['W', 'vegas_win_total', 'abs_line_acc']
compare_lines2['pdiff_line_accuracy'] = (compare_lines2['W'] - compare_lines2['vegas_win_total'])
compare_lines2
```

Out[329]:

| | W | vegas_win_total | abs_line_accuracy | pdiff_line_accuracy |
|------------|------|-----------------|-------------------|---------------------|
| yearID | | | | |
| 2014-01-01 | 2430 | 2454.500 | -24.500 | -0.010 |
| 2015-01-01 | 2429 | 2436.000 | -7.000 | -0.003 |
| 2016-01-01 | 2427 | 2436.500 | -9.500 | -0.004 |
| 2017-01-01 | 2430 | 2447.500 | -17.500 | -0.007 |
| 2018-01-01 | 2431 | 2433.000 | -2.000 | -0.001 |
| 2019-01-01 | 2429 | 2431.500 | -2.500 | -0.001 |

```
In [330]: # Vegas accuracy by team cumulative 2014 - 2019
compare_lines3 = vegas_df.groupby(['franchID'])['W', 'vegas_win_total', 'abs_line_a
compare_lines3['pdiff_line_accuracy'] = (compare_lines3['W'] - compare_lines3['ve
compare_lines3.sort_values('pdiff_line_accuracy', ascending=False)
```

Out[330]:

| | W | vegas_win_total | abs_line_accuracy | pdiff_line_accuracy |
|----------|-----|-----------------|-------------------|---------------------|
| franchID | | | | |
| HOU | 551 | 508.000 | 43.000 | 0.085 |
| MIL | 494 | 464.000 | 30.000 | 0.065 |
| NYN | 549 | 524.000 | 25.000 | 0.048 |
| CHC | 544 | 524.000 | 20.000 | 0.038 |
| OAK | 494 | 476.500 | 17.500 | 0.037 |
| CLE | 546 | 527.500 | 18.500 | 0.035 |
| LAD | 579 | 560.500 | 18.500 | 0.033 |
| ATL | 473 | 459.000 | 14.000 | 0.031 |
| MIN | 476 | 462.000 | 14.000 | 0.030 |
| TBD | 491 | 478.000 | 13.000 | 0.027 |
| STL | 538 | 524.000 | 14.000 | 0.027 |
| WSN | 546 | 546.000 | 0.000 | 0.000 |
| PIT | 490 | 492.500 | -2.500 | -0.005 |
| ARI | 472 | 475.000 | -3.000 | -0.006 |
| KCR | 462 | 469.000 | -7.000 | -0.015 |
| COL | 458 | 465.000 | -7.000 | -0.015 |
| NYM | 489 | 497.000 | -8.000 | -0.016 |
| SEA | 484 | 492.000 | -8.000 | -0.016 |
| PHI | 434 | 441.500 | -7.500 | -0.017 |
| BOS | 527 | 540.500 | -13.500 | -0.025 |
| FLA | 424 | 437.000 | -13.000 | -0.030 |
| ANA | 489 | 504.000 | -15.000 | -0.030 |
| TOR | 481 | 496.000 | -15.000 | -0.030 |
| TEX | 473 | 489.000 | -16.000 | -0.033 |
| BAL | 442 | 458.000 | -16.000 | -0.035 |
| CHW | 428 | 451.500 | -23.500 | -0.052 |
| SFG | 473 | 499.000 | -26.000 | -0.052 |
| SDP | 426 | 452.000 | -26.000 | -0.058 |
| CIN | 418 | 453.000 | -35.000 | -0.077 |
| DET | 425 | 473.500 | -48.500 | -0.102 |

As you can see from the above, the bookmakers are EXTREMELY accurate!

With that said, there is value in looking into the teams with the largest differences and analyzing how/why the bookmakers are off on their predictions for these teams.

```
In [291]: print(df_teams.shape)
          print(vegas_df.shape)
```

```
(600, 32)
(180, 37)
```

```
In [292]: import pickle
          df_teams.to_pickle('df_teams.pkl')
          vegas_df.to_pickle('vegas_df.pkl')
```

Single Game Results

With our knowledge of season total wins, it is now time to dig into single game results. From the lahman's baseball database, we have every game's stats and results from 2014 - 2019.

```
In [293]: df_gamelogs = pd.read_csv('data/gamelogs.csv')
          df_gamelogs.head()
```

Out[293]:

| | date | dayofweek | visiting_team | v_game_number | home_team | h_game_number | v_score | h_ |
|---|----------|-----------|---------------|---------------|-----------|---------------|---------|----|
| 0 | 20190320 | Wed | SEA | 1 | OAK | 1 | 9 | |
| 1 | 20190321 | Thu | SEA | 2 | OAK | 2 | 5 | |
| 2 | 20190328 | Thu | PIT | 1 | CIN | 1 | 3 | |
| 3 | 20190328 | Thu | ARI | 1 | LAD | 1 | 5 | |
| 4 | 20190328 | Thu | COL | 1 | FLA | 1 | 6 | |

```
In [297]: # there were 16 missing values to fill  
df_gamelogs['LOG_minutes'].fillna(method='backfill', inplace=True)  
df_gamelogs.isna().sum()
```

```
Out[297]: date                0  
dayofweek                   0  
visiting_team               0  
v_game_number              0  
home_team                   0  
h_game_number              0  
v_score                     0  
h_score                     0  
LOG_inouts                  0  
LOG_minutes                 0  
v_AB                        0  
v_H                         0  
v_2B                        0  
v_3B                        0  
v_HR                        0  
v_RBI                       0  
v_SF                        0  
v_HBP                       0  
v_BB                        0  
v                            0
```

```
In [298]: df_gamelogs['date'] = pd.to_datetime(df_gamelogs['date'], format='%Y%m%d')
df_gamelogs.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 14577 entries, 0 to 14576
Data columns (total 46 columns):
date                14577 non-null datetime64[ns]
dayofweek           14577 non-null object
visiting_team       14577 non-null object
v_game_number       14577 non-null int64
home_team           14577 non-null object
h_game_number       14577 non-null int64
v_score             14577 non-null int64
h_score             14577 non-null int64
LOG_inouts          14577 non-null int64
LOG_minutes         14577 non-null float64
v_AB                14577 non-null int64
v_H                 14577 non-null int64
v_2B                14577 non-null int64
v_3B                14577 non-null object
v_HR                14577 non-null int64
v_RBI               14577 non-null int64
v_SF                14577 non-null int64
v_HBP               14577 non-null int64
v_BB                14577 non-null int64
v_K                 14577 non-null int64
v_SB                14577 non-null int64
v_GIDP              14577 non-null int64
v_LOB               14577 non-null int64
v_pitchers_used     14577 non-null int64
v_ER                14577 non-null int64
v_def_assists       14577 non-null int64
v_errors_def        14577 non-null int64
v_doubleplays_def   14577 non-null int64
h_AB                14577 non-null int64
h_H                 14577 non-null int64
h_2B                14577 non-null int64
h_3B                14577 non-null int64
h_HR                14577 non-null int64
h_RBI               14577 non-null int64
h_SF                14577 non-null int64
h_HBP               14577 non-null int64
h_BB                14577 non-null int64
h_K                 14577 non-null int64
h_SB                14577 non-null int64
h_GIDP              14577 non-null int64
h_LOB               14577 non-null int64
h_pitchers_used     14577 non-null int64
h_ER                14577 non-null int64
h_def_assists       14577 non-null int64
h_errors_def        14577 non-null int64
h_doubleplays_def   14577 non-null int64
dtypes: datetime64[ns](1), float64(1), int64(40), object(4)
memory usage: 5.1+ MB
```

```
In [299]: df_gamelogs.describe()
```

```
Out[299]:
```

| | v_game_number | h_game_number | v_score | h_score | LOG_inouts | LOG_minutes | v_ |
|--------------|---------------|---------------|-----------|-----------|------------|-------------|---------|
| count | 14577.000 | 14577.000 | 14577.000 | 14577.000 | 14577.000 | 14577.000 | 14577.0 |
| mean | 81.484 | 81.483 | 4.392 | 4.514 | 53.611 | 186.157 | 68.8 |
| std | 46.756 | 46.758 | 3.176 | 3.122 | 5.108 | 28.071 | 1058.6 |
| min | 1.000 | 1.000 | 0.000 | 0.000 | 27.000 | 75.000 | 17.0 |
| 25% | 41.000 | 41.000 | 2.000 | 2.000 | 51.000 | 168.000 | 32.0 |
| 50% | 82.000 | 81.000 | 4.000 | 4.000 | 54.000 | 183.000 | 34.0 |
| 75% | 122.000 | 122.000 | 6.000 | 6.000 | 54.000 | 200.000 | 37.0 |
| max | 163.000 | 163.000 | 24.000 | 25.000 | 114.000 | 413.000 | 41567.0 |

```
In [300]: df_gamelogs.set_index('date', inplace=True)
```

```
In [301]: df_gamelogs = df_gamelogs.sort_index()
df_gamelogs
```

```
Out[301]:
```

| | dayofweek | visiting_team | v_game_number | home_team | h_game_number | v_score | h_score |
|------------|-----------|---------------|---------------|-----------|---------------|---------|---------|
| date | | | | | | | |
| 2014-03-22 | Sat | LAD | 1 | ARI | 1 | 3 | 1 |
| 2014-03-23 | Sun | LAD | 2 | ARI | 2 | 7 | 5 |
| 2014-03-30 | Sun | LAD | 3 | SDP | 1 | 1 | 3 |
| 2014-03-31 | Mon | ATL | 1 | MIL | 1 | 0 | 2 |
| 2014-03-31 | Mon | WSN | 1 | NYM | 1 | 9 | 7 |
| ... | ... | ... | ... | ... | ... | ... | ... |
| 2019-09-29 | Sun | BAL | 162 | BOS | 162 | 4 | 5 |
| 2019-09-29 | Sun | DET | 161 | CHW | 161 | 3 | 5 |
| 2019-09-29 | Sun | MIN | 162 | KCR | 162 | 4 | 5 |
| 2019-09-29 | Sun | NYN | 162 | TEX | 162 | 1 | 6 |
| 2019-09-29 | Sun | HOU | 162 | ANA | 162 | 8 | 5 |

14577 rows × 45 columns

```
In [302]: df_gamelogs.to_pickle('df_gamelogs.pkl')
```


Rolling Statistics:

```
In [303]: df = pd.read_pickle('df_gamelogs.pkl')
df
```

Out[303]:

| | dayofweek | visiting_team | v_game_number | home_team | h_game_number | v_score | h_score |
|------------|-----------|---------------|---------------|-----------|---------------|---------|---------|
| date | | | | | | | |
| 2014-03-22 | Sat | LAD | 1 | ARI | 1 | 3 | 1 |
| 2014-03-23 | Sun | LAD | 2 | ARI | 2 | 7 | 5 |
| 2014-03-30 | Sun | LAD | 3 | SDP | 1 | 1 | 3 |
| 2014-03-31 | Mon | ATL | 1 | MIL | 1 | 0 | 2 |
| 2014-03-31 | Mon | WSN | 1 | NYM | 1 | 9 | 7 |
| ... | ... | ... | ... | ... | ... | ... | ... |
| 2019-09-29 | Sun | BAL | 162 | BOS | 162 | 4 | 5 |
| 2019-09-29 | Sun | DET | 161 | CHW | 161 | 3 | 5 |
| 2019-09-29 | Sun | MIN | 162 | KCR | 162 | 4 | 5 |
| 2019-09-29 | Sun | NYN | 162 | TEX | 162 | 1 | 6 |
| 2019-09-29 | Sun | HOU | 162 | ANA | 162 | 8 | 5 |

14577 rows × 45 columns



```
In [304]: df.reset_index(inplace=True)
```

```
In [305]: teams = list(df['visiting_team'].value_counts().index)
          teams
```

```
Out[305]: ['MIL',
           'COL',
           'SFG',
           'LAD',
           'NYY',
           'STL',
           'HOU',
           'TOR',
           'CHC',
           'TEX',
           'WSN',
           'BAL',
           'ARI',
           'CHW',
           'ANA',
           'MIN',
           'CIN',
           'PIT',
           'KCR',
           'NOC']
```

```
In [306]: # create blank df with features needed for analysis
          new_df = pd.DataFrame(columns = ['date', 'team', 'team_score', 'opponent_score',
          'Hits', 'RBI', 'doubles', 'strikeouts', 'def_assists', 'def_errors', 'def_dou
          'pitchers_used'])
```

```
In [43]: df.columns
```

```
Out[43]: Index(['date', 'dayofweek', 'visiting_team', 'v_game_number', 'home_team',
               'h_game_number', 'v_score', 'h_score', 'LOG_inouts', 'LOG_minutes',
               'v_AB', 'v_H', 'v_2B', 'v_3B', 'v_HR', 'v_RBI', 'v_SF', 'v_HBP', 'v_BB',
               'v_K', 'v_SB', 'v_GIDP', 'v_LOB', 'v_pitchers_used', 'v_ER',
               'v_def_assists', 'v_errors_def', 'v_doubleplays_def', 'h_AB', 'h_H',
               'h_2B', 'h_3B', 'h_HR', 'h_RBI', 'h_SF', 'h_HBP', 'h_BB', 'h_K', 'h_SB',
               'h_GIDP', 'h_LOB', 'h_pitchers_used', 'h_ER', 'h_def_assists',
               'h_errors_def', 'h_doubleplays_def'],
               dtype='object')
```



```
In [307]: v_temp = df.loc[df['visiting_team'] == 'SEA',['date', 'visiting_team', 'v_score',
                                                    'v_2B','v_K','v_def_ass',
                                                    'v_H': 'Hits','v_RBI':
                                                    'v_errors_def': 'def_er

v_temp
```

Out[307]:

| | date | team | team_score | opponent_score | Hits | RBI | doubles | strikeouts | def_assists | def_ |
|--------------|------------|------|------------|----------------|------|-----|---------|------------|-------------|------|
| 13 | 2014-03-31 | SEA | 10 | 3 | 11 | 10 | 4 | 11 | 5 | |
| 21 | 2014-04-01 | SEA | 8 | 3 | 10 | 8 | 3 | 12 | 14 | |
| 38 | 2014-04-02 | SEA | 8 | 2 | 13 | 8 | 2 | 9 | 11 | |
| 45 | 2014-04-03 | SEA | 2 | 3 | 6 | 2 | 1 | 9 | 11 | |
| 74 | 2014-04-05 | SEA | 3 | 1 | 7 | 3 | 2 | 9 | 7 | |
| ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... |
| 14425 | 2019-09-18 | SEA | 4 | 1 | 7 | 4 | 2 | 10 | 12 | |
| 14436 | 2019-09-19 | SEA | 6 | 5 | 13 | 5 | 1 | 12 | 13 | |
| 14447 | 2019-09-20 | SEA | 3 | 5 | 4 | 3 | 1 | 5 | 9 | |
| 14469 | 2019-09-21 | SEA | 7 | 6 | 11 | 6 | 1 | 10 | 16 | |
| 14477 | 2019-09-22 | SEA | 1 | 2 | 7 | 1 | 3 | 5 | 11 | |

486 rows × 12 columns

```
In [308]: h_temp = df.loc[df['home_team'] == 'SEA',['date', 'home_team', 'v_score', 'h_score',
                                                    'h_2B','h_K','h_def_assists',
                                                    'h_H': 'Hits','h_RBI': 'RBI',
                                                    'h_errors_def': 'def_errors']]
h_temp
```

Out[308]:

| | date | team | opponent_score | team_score | Hits | RBI | doubles | strikeouts | def_assists | def_errors |
|-------|------------|------|----------------|------------|------|-----|---------|------------|-------------|------------|
| 107 | 2014-04-08 | SEA | 3 | 5 | 8 | 5 | 2 | 9 | 10 | |
| 123 | 2014-04-09 | SEA | 2 | 0 | 1 | 0 | 0 | 9 | 9 | |
| 149 | 2014-04-11 | SEA | 4 | 6 | 11 | 4 | 2 | 12 | 3 | |
| 164 | 2014-04-12 | SEA | 3 | 1 | 7 | 1 | 1 | 11 | 9 | |
| 178 | 2014-04-13 | SEA | 3 | 0 | 3 | 0 | 0 | 11 | 6 | |
| ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... |
| 14509 | 2019-09-25 | SEA | 3 | 0 | 2 | 0 | 0 | 10 | 8 | |
| 14522 | 2019-09-26 | SEA | 3 | 1 | 4 | 1 | 1 | 10 | 9 | |
| 14534 | 2019-09-27 | SEA | 3 | 4 | 8 | 4 | 3 | 9 | 7 | |
| 14548 | 2019-09-28 | SEA | 1 | 0 | 6 | 0 | 0 | 6 | 11 | |
| 14562 | 2019-09-29 | SEA | 1 | 3 | 7 | 3 | 2 | 7 | 8 | |

486 rows × 12 columns

```
In [309]: #Loop through each team, create temp df for when they are home and away, rename columns
for team in teams:
    v_temp = df.loc[df['visiting_team'] == team,['date', 'visiting_team', 'v_score', 'v_2B', 'v_K', 'v_def_assists', 'v_errors',
                                                  'v_H': 'Hits', 'v_RBI': 'RBI', 'v_doubles': 'doubles',
                                                  'v_errors_def': 'def_errors', 'v_doubles_def': 'def_doubles']]
    h_temp = df.loc[df['home_team'] == team,['date', 'home_team', 'v_score', 'h_score', 'h_2B', 'h_K', 'h_def_assists', 'h_errors',
                                              'h_H': 'Hits', 'h_RBI': 'RBI', 'h_2B': '2B', 'h_K': 'K',
                                              'h_errors_def': 'def_errors', 'h_doubles': 'doubles']]
    new_df = pd.concat([new_df, v_temp, h_temp])
```

```
In [310]: new_df.shape
```

```
Out[310]: (29154, 12)
```

```
In [311]: new_df.head()
```

```
Out[311]:
```

| | Hits | RBI | date | def_assists | def_doubleplays | def_errors | doubles | opponent_score | pitchers |
|------------|------|-----|------------|-------------|-----------------|------------|---------|----------------|----------|
| 62 | 12 | 6 | 2014-04-04 | 7 | 2 | 1 | 4 | 2 | |
| 71 | 19 | 7 | 2014-04-05 | 14 | 2 | 2 | 6 | 6 | |
| 90 | 9 | 3 | 2014-04-06 | 10 | 0 | 1 | 2 | 0 | |
| 104 | 15 | 9 | 2014-04-08 | 4 | 0 | 0 | 5 | 4 | |
| 118 | 13 | 8 | 2014-04-09 | 11 | 1 | 1 | 3 | 4 | |



```
In [312]: new_df.sort_values('date', inplace= True)
```

```
In [313]: new_df = new_df[['date', 'team', 'team_score', 'opponent_score',  
                          'Hits','RBI','doubles','strikeouts', 'def_assists', 'def_errors', 'def_dou  
                          'pitchers_used']]
```

In [314]: new_df

Out[314]:

| | date | team | team_score | opponent_score | Hits | RBI | doubles | strikeouts | def_assists | def_ |
|-------|------------|------|------------|----------------|------|-----|---------|------------|-------------|------|
| 0 | 2014-03-22 | ARI | 1 | 3 | 5 | 1 | 1 | 10 | 10 | |
| 0 | 2014-03-22 | LAD | 3 | 1 | 5 | 3 | 2 | 11 | 13 | |
| 1 | 2014-03-23 | LAD | 7 | 5 | 13 | 6 | 3 | 7 | 4 | |
| 1 | 2014-03-23 | ARI | 5 | 7 | 8 | 5 | 0 | 8 | 15 | |
| 2 | 2014-03-30 | SDP | 3 | 1 | 5 | 3 | 0 | 10 | 10 | |
| ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | |
| 14573 | 2019-09-29 | CHW | 5 | 3 | 4 | 4 | 2 | 12 | 15 | |
| 14566 | 2019-09-29 | PHI | 3 | 4 | 13 | 3 | 1 | 10 | 7 | |
| 14563 | 2019-09-29 | ARI | 1 | 0 | 7 | 1 | 2 | 8 | 14 | |
| 14569 | 2019-09-29 | STL | 9 | 0 | 9 | 8 | 0 | 7 | 11 | |
| 14568 | 2019-09-29 | LAD | 9 | 0 | 12 | 9 | 1 | 8 | 10 | |

29154 rows × 12 columns



```
In [315]: # calculate rolling statistics
new_df['Rolling3_runs_scored'] = new_df.groupby('team')['team_score'].transform(lambda x: x.rolling(3).sum())
new_df['Rolling3_runs_against'] = new_df.groupby('team')['opponent_score'].transform(lambda x: x.rolling(3).sum())
new_df['Rolling3_run_diff'] = new_df['Rolling3_runs_scored'] - new_df['Rolling3_runs_against']
new_df['Rolling3_hits'] = new_df.groupby('team')['Hits'].transform(lambda x: x.rolling(3).sum())
new_df['Rolling3_RBI'] = new_df.groupby('team')['RBI'].transform(lambda x: x.rolling(3).sum())
new_df['Rolling3_doubles'] = new_df.groupby('team')['doubles'].transform(lambda x: x.rolling(3).sum())
new_df['Rolling3_strikeouts'] = new_df.groupby('team')['strikeouts'].transform(lambda x: x.rolling(3).sum())
new_df['Rolling3_def_assists'] = new_df.groupby('team')['def_assists'].transform(lambda x: x.rolling(3).sum())
new_df['Rolling3_def_errors'] = new_df.groupby('team')['def_errors'].transform(lambda x: x.rolling(3).sum())
new_df['Rolling3_doubleplays'] = new_df.groupby('team')['def_doubleplays'].transform(lambda x: x.rolling(3).sum())
new_df['Rolling3_pitchers_used'] = new_df.groupby('team')['pitchers_used'].transform(lambda x: x.rolling(3).sum())
```

```
In [316]: # create target column that defines the game result as a 1 or 0
new_df['game_result'] = np.where(new_df['team_score'] - new_df['opponent_score'] > 0, 1, 0)
```

```
In [317]: new_df.fillna(method='backfill', inplace=True)
          new_df.isna().sum()
```

```
Out[317]: date                0
          team                0
          team_score          0
          opponent_score      0
          Hits                0
          RBI                 0
          doubles             0
          strikeouts          0
          def_assists         0
          def_errors          0
          def_doubleplays     0
          pitchers_used       0
          Rolling3_runs_scored 0
          Rolling3_runs_against 0
          Rolling3_run_diff    0
          Rolling3_hits       0
          Rolling3_RBI        0
          Rolling3_doubles     0
          Rolling3_strikeouts  0
          Rolling3_def_assists 0
          Rolling3_def_errors  0
          Rolling3_doubleplays 0
          Rolling3_pitchers_used 0
          game_result          0
          dtype: int64
```

```
In [318]: new_df
```

Out[318]:

| | date | team | team_score | opponent_score | Hits | RBI | doubles | strikeouts | def_assists | def_ |
|-------|------------|------|------------|----------------|------|-----|---------|------------|-------------|------|
| 0 | 2014-03-22 | ARI | 1 | 3 | 5 | 1 | 1 | 10 | 10 | |
| 0 | 2014-03-22 | LAD | 3 | 1 | 5 | 3 | 2 | 11 | 13 | |
| 1 | 2014-03-23 | LAD | 7 | 5 | 13 | 6 | 3 | 7 | 4 | |
| 1 | 2014-03-23 | ARI | 5 | 7 | 8 | 5 | 0 | 8 | 15 | |
| 2 | 2014-03-30 | SDP | 3 | 1 | 5 | 3 | 0 | 10 | 10 | |
| ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... |
| 14573 | 2019-09-29 | CHW | 5 | 3 | 4 | 4 | 2 | 12 | 15 | |
| 14566 | 2019-09-29 | PHI | 3 | 4 | 13 | 3 | 1 | 10 | 7 | |
| 14563 | 2019-09-29 | ARI | 1 | 0 | 7 | 1 | 2 | 8 | 14 | |
| 14569 | 2019-09-29 | STL | 9 | 0 | 9 | 8 | 0 | 7 | 11 | |
| 14568 | 2019-09-29 | LAD | 9 | 0 | 12 | 9 | 1 | 8 | 10 | |

29154 rows × 24 columns



```
In [319]: check = new_df.loc[(new_df['team'] == 'NYM') | (new_df['team'] == 'NYM')]
check.head(20)
```

Out[319]:

| | date | team | team_score | opponent_score | Hits | RBI | doubles | strikeouts | def_assists | def_er |
|------------|------------|------|------------|----------------|------|-----|---------|------------|-------------|--------|
| 4 | 2014-03-31 | NYM | 7 | 9 | 7 | 7 | 0 | 18 | 11 | |
| 26 | 2014-04-02 | NYM | 1 | 5 | 3 | 1 | 1 | 13 | 6 | |
| 42 | 2014-04-03 | NYM | 2 | 8 | 7 | 2 | 3 | 8 | 11 | |
| 53 | 2014-04-04 | NYM | 4 | 3 | 6 | 4 | 1 | 4 | 11 | |
| 65 | 2014-04-05 | NYM | 6 | 3 | 6 | 6 | 1 | 10 | 6 | |
| 80 | 2014-04-06 | NYM | 1 | 2 | 4 | 1 | 2 | 8 | 9 | |
| 108 | 2014-04-08 | NYM | 4 | 0 | 9 | 3 | 1 | 9 | 8 | |
| 129 | 2014-04-09 | NYM | 3 | 4 | 6 | 3 | 1 | 9 | 6 | |
| 134 | 2014-04-10 | NYM | 6 | 4 | 9 | 6 | 1 | 9 | 12 | |
| 153 | 2014-04-11 | NYM | 4 | 5 | 10 | 3 | 3 | 9 | 11 | |
| 167 | 2014-04-12 | NYM | 7 | 6 | 11 | 7 | 0 | 14 | 19 | |
| 183 | 2014-04-13 | NYM | 2 | 14 | 6 | 1 | 1 | 11 | 9 | |
| 189 | 2014-04-14 | NYM | 7 | 3 | 13 | 7 | 2 | 10 | 14 | |
| 204 | 2014-04-15 | NYM | 9 | 0 | 12 | 8 | 1 | 10 | 6 | |
| 211 | 2014-04-16 | NYM | 5 | 2 | 10 | 5 | 2 | 8 | 10 | |
| 235 | 2014-04-18 | NYM | 0 | 6 | 1 | 0 | 0 | 8 | 11 | |
| 250 | 2014-04-19 | NYM | 5 | 7 | 12 | 5 | 3 | 8 | 8 | |
| 264 | 2014-04-20 | NYM | 4 | 3 | 9 | 4 | 0 | 11 | 13 | |
| 279 | 2014-04-21 | NYM | 2 | 0 | 7 | 2 | 1 | 8 | 11 | |
| 292 | 2014-04-22 | NYM | 0 | 3 | 4 | 0 | 0 | 4 | 8 | |

```
In [320]: # this is the mets record from 2014 - 2019 to confirm data is still in tact  
check.game_result.value_counts()
```

```
Out[320]: 1    489  
         0    483  
         Name: game_result, dtype: int64
```

```
In [321]: new_df.to_pickle('new_df.pkl')
```

```
In [ ]:
```

```
In [ ]:
```